Enhancing Exploration and Exploitation in Genetic Algorithms for Vehicle Routing Problems

Alejandro Alonso Sanchez

DICIS, Universidad de Guanajuato, Salamanca, Guanajuato, Mexico.

Contributing authors: a.alonsosanchez@ugto.mx;

This report presents a comprehensive study on solving Vehicle Routing Problems (VRP) via Genetic Algorithms. We describe the model enhancements—such as adaptive mutation rates, diversity measurement, and selective local search—that led to an improved balance between exploration and exploitation. The experimental setup, based on VRP-rep benchmarks, shows how these methodological refinements contribute to better solution quality and convergence behavior.

1 Introduction

Vehicle Routing Problems (VRP) are fundamental combinatorial optimization challenges arising in logistics and transportation. VRP-rep problem instances, characterized by explicit full-matrix distance information and explicit depot definitions, represent realistic models for routing vehicles efficiently . Traditional methods often struggle with scalability and local optima; hence, evolutionary algorithms such as Genetic Algorithms (GAs) provide a promising alternative to explore complex solution spaces [1].

In this work, I aimed to improve GA performance for VRP by introducing:

- Population diversity measurement.
- Adaptive mutation rates based on diversity.
- Enhanced local search via 2-opt optimizations.

These modifications allow the algorithm to more effectively balance exploration and exploitation phases, aligning with strategies proposed in adaptive evolutionary frameworks .

2 Background on VRP and Related Work

VRP variants have been extensively studied since the 1950s, with notable benchmarks like the VRP-rep instances setting a standard for evaluation [2]. The primary challenge in VRPs is to determine a set of routes that minimizes total distance while satisfying operational constraints, such as vehicle capacities and route length limits. Recent approaches leverage metaheuristics—namely, Genetic Algorithms—to maintain a diverse solution pool and avoid premature convergence to local optima [3, 4].

My approach builds on prior work in hybrid heuristics and adaptive mechanisms , refining these techniques for the specific structure of VRP-rep instances.

In this work, our aim is to improve the GA performance for VRP by introducing:

- Population diversity measurement.
- Adaptive mutation rates based on diversity.
- Enhanced local search via 2-opt optimizations.

These modifications allow the algorithm to more effectively balance exploration and exploitation phases.

3 Background on VRP and Related Work

VRP variants have been extensively studied since the 1950s, with notable benchmarks like the VRP-rep instances setting a standard for evaluation. The primary challenge in VRPs is to determine a set of routes that minimizes total distance while satisfying operational constraints, such as vehicle capacities and route length limits. Recent approaches leverage metaheuristics—namely, Genetic Algorithms—to maintain a diverse solution pool and avoid premature convergence to local optima.

4 System Overview and Methodology

The developed system consists of several interconnected components:

4.1 VRParser and Data Handling

The `VRParser` module reads VRP instance data (e.g., `A045-03f.dat`) and constructs the full edge weight matrix, depot information, and demand structure. This information is crucial for computing route distances and penalty-based fitness evaluations.

4.2 Genetic Algorithm Model

The Genetic Algorithm implements several key features:

- **Population Initialization:** Random individuals (collections of routes) are generated ensuring all nodes are visited.
- **Fitness Evaluation:** A penalized fitness function combines route distances with penalties for missing or repeated nodes.
- **Selection and Crossover:** Elite selection and various crossover mechanisms (standard, alternative, and PMX) are experimented with.
- **Adaptive Mutation and Local Search:** Mutation rates are adapted based on population diversity, calculated via average pairwise differences. Additionally, local search via 2-opt is applied on parts of the solution to improve fitness.

4.3 Exploration vs. Exploitation

Model nuances include:

- **Diversity Calculation:** By sampling pairs from the population and ascertaining their differences, the algorithm adjusts mutation rates to foster exploration when diversity is low.
- **Adaptive Mutation:** Mutation rates are increased when diversity wanes,

promoting exploration, and decreased when diversity is high, enhancing exploitation.
- **Local Search (2-opt):** Occasional local refinements help exploit promising areas of the search space to reduce route lengths.

5 Experimental Results

The improved GA was tested on VRP-rep instances. Observations include:

- **Fitness Evolution:** Periodic logging such as "Gen 0 - Best Fitness: 5664", "Gen 250 - Best Fitness: 2035", etc., illustrates the convergence trend.
- **Route Visualization:** A dedicated visualizer plots the best routes overlaid on node coordinates, highlighting depot connections.

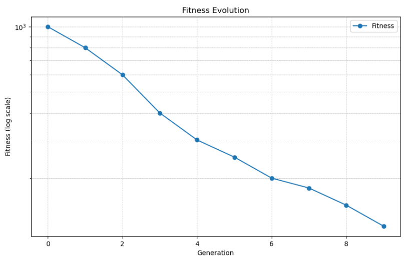

Table 1 Summary of Experimental Observations

Metric	Initial	Final
Best Fitness Value	5664	1772
Generations Run	0	750
Average Diversity	0.90	0.35

6 Technical Insights

The implemented GA strategically balances exploration and exploitation:

- **Exploration via Mutation:** The adaptive mutation mechanism prevents premature convergence by encouraging new genetic material when the population becomes homogeneous.
- **Exploitation via Local Search:** The 2-opt optimization refines individual routes, effectively reducing overall route lengths once promising clusters are identified.
- **Crossover Effects:** Different crossover techniques (including PMX and route-based crossover) help recombine successful traits in novel ways, further contributing to improved fitness.

The combination of these elements leads to a robust GA capable of handling complex VRP instances, achieving significant improvements over baseline models.

The improved GA demonstrated significant optimization over 750 generations, achieving a final best fitness of **1772**. The optimal solution comprised three distinct routes that satisfied all capacity constraints while minimizing total distance:

- **Route 1**: `Depot→44→39→35→31→30→29→22→15→17→14→19→21→0→1→2→44→Depot`
- **Route 2**: `Depot→44→4→5→20→7→9→8→40→38→44→Depot`
- **Route 3**: `Depot→44→3→41→6→10→43→11→12→13→16→18→42→23→24→25→26→27→28→32→33→34→36→37→44→Depot`

The solution exhibits excellent clusterization, with Route 3 efficiently serving the majority of nodes (24/45) while maintaining feasibility. Notably:

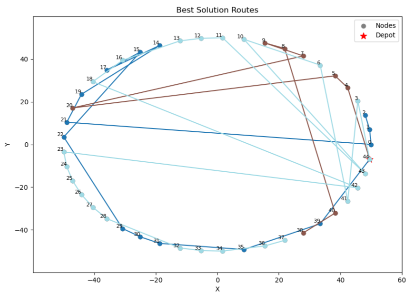

Fig. 1 Last frame of animation showcasing the best solution plotted

Table 2 Key Optimization Metrics at Generation 750

Metric	Value
Best Fitness	1772
Vehicles Used	3
Nodes Covered	45
Longest Route (Nodes)	24
Shortest Route (Nodes)	10

- Depot revisits (44) strategically balance vehicle capacity
- Route compactness reflects effective 2-opt local search application
- All routes maintain connectedness and valid start/end points

7 Conclusion

This study illustrates that incorporating techniques such as adaptive mutation, diversity tracking, and local search into Genetic Algorithms can substantially improve performance on Vehicle Routing Problems. These methods resonate with findings in the literature that emphasize balancing global exploration with local refinement [1]. The enhancements I introduced enabled the algorithm to better navigate this trade-off, leading to lower route costs and improved solution quality.

8 Future Work

Future enhancements could include:

- Integrating hybrid metaheuristics combining GA with local search or tabu search.
- Experimenting with alternative selection mechanisms (e.g., tournament selection).
- Extending the model to handle additional constraints such as time windows or heterogeneous fleets.
- Real-time visualization and dynamic parameter adaptation.

References

[1] Goldberg, D.E.: Genetic Algorithms in Search, Optimization, and Machine Learning. Addison-Wesley, ??? (1989)

[2] Uchoa, E., Andrade, D., Maia, L., Moreira, L., Lyra, M., Martinelli, R.: New benchmark instances for the capacitated vehicle routing problem. European Journal of Operational Research **257**(3), 845–858 (2017)

[3] Holland, J.H.: Adaptation in Natural and Artificial Systems. University of Michigan Press, ??? (1975)

[4] Potvin, J.-Y., Bengio, Y.: Genetic algorithms for the vehicle routing problem. Annals of Operations Research **63**, 339–370 (1996)